

Project: No Show Appointments Data Analysis

Table of Contents

- [Introduction](#)
- [Data Wrangling](#)
- [Data Cleaning Procedure](#)
- [Cleaned Data](#)
- [Exploratory Data Analysis](#)
 - [Research Question 1: Does Gender affect the rate of show up?](#)
 - [Research Question 2: Does the Bosnia Familia Scholarship affect the rate of show up?](#)
 - [Research Question 3: What is the effect of disability on show up?](#)
 - [Research Question 4: How was the bosnia scholarship distributed among patients?](#)
- [Conclusions](#)

Introduction

This dataset collects information from about 100k medical appointments in Brazil and is focused on the question of whether or not patients show up for their appointment. A number of characteristics about the patient are included in each row. I will be analyzing this dataset to answer the following questions.

- Does Gender affect the rate of show up?
- Does the Bosnia Familia Scholarship affect the rate of show up?
- What is the effect of disability on show up?
- How was the bosnia scholarship distributed among patients?

Import all necessary libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import re

#To display entire dataset
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

Data Wrangling

Data Description (source: <https://www.kaggle.com/joniarroba/noshowappointments> (<https://www.kaggle.com/joniarroba/noshowappointments>))

- 01 - Patient_Id.....Identification of a patient
- 02 - AppointmentID.....Identification of each appointment
- 03 - Gender.....(Male or Female).
- 04 - ScheduledDay.....The day someone called or registered the appointment, this is before appointment of course.
- 05 - AppointmentDay.....The day of the actual appointment, when they have to visit the doctor.
- 06 - Age.....How old is the patient.
- 07 - Neighbourhood.....Where the appointment takes place.
- 08 - Scholarship.....True of False
- 09 - Hipertension.....True of False
- 10 - Diabetes.....True of False
- 11 - Alcoholism.....True of False
- 12 - Handcap.....True of False
- 13 - SMS_received.....1 or more messages sent to the patient.
- 14 - No-show.....True of False

Load data and check first few rows and columns.

In [2]: `df = pd.read_csv('C:/Users/user/Desktop/1.coding_datascience/UDACITY NANO DEGREE
df.head()`

Out[2]:

	PatientId	AppointmentID	Gender	ScheduledDay	AppointmentDay	Age	Neighbourhood	Sc
0	2.987250e+13	5642903	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	
1	5.589978e+14	5642503	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	
2	4.262962e+12	5642549	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA	
3	8.679512e+11	5642828	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI	
4	8.841186e+12	5642494	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	

check number of observations and variables (we have 110527 observations and 14 variables)

In [3]: `print(df.shape)`
(110527, 14)

check for missing data (there are no missing data)

```
In [4]: df.isnull().sum()
```

```
Out[4]: PatientId      0
AppointmentID    0
Gender           0
ScheduledDay     0
AppointmentDay   0
Age             0
Neighbourhood    0
Scholarship     0
Hipertension     0
Diabetes         0
Alcoholism      0
Handcap         0
SMS_received    0
No-show         0
dtype: int64
```

check for duplicates entries (there are no duplicates entries)

```
In [5]: df.duplicated().sum()
```

```
Out[5]: 0
```

check unique values in each column

```
In [6]: df.nunique()
```

```
Out[6]: PatientId      62299
AppointmentID    110527
Gender           2
ScheduledDay     103549
AppointmentDay   27
Age             104
Neighbourhood    81
Scholarship     2
Hipertension     2
Diabetes         2
Alcoholism      2
Handcap         5
SMS_received    2
No-show         2
dtype: int64
```

Check data types

In [7]: `df.dtypes`

```
Out[7]: PatientId      float64
AppointmentID    int64
Gender           object
ScheduledDay     object
AppointmentDay   object
Age             int64
Neighbourhood    object
Scholarship      int64
Hypertension     int64
Diabetes         int64
Alcoholism       int64
Handcap          int64
SMS_received     int64
No-show         object
dtype: object
```

Summary Statistics

In [8]: `df.describe()`

```
Out[8]:
```

	PatientId	AppointmentID	Age	Scholarship	Hypertension	Diabetes
count	1.105270e+05	1.105270e+05	110527.000000	110527.000000	110527.000000	110527.000000
mean	1.474963e+14	5.675305e+06	37.088874	0.098266	0.197246	0.071865
std	2.560949e+14	7.129575e+04	23.110205	0.297675	0.397921	0.258265
min	3.921784e+04	5.030230e+06	-1.000000	0.000000	0.000000	0.000000
25%	4.172614e+12	5.640286e+06	18.000000	0.000000	0.000000	0.000000
50%	3.173184e+13	5.680573e+06	37.000000	0.000000	0.000000	0.000000
75%	9.439172e+13	5.725524e+06	55.000000	0.000000	0.000000	0.000000
max	9.999816e+14	5.790484e+06	115.000000	1.000000	1.000000	1.000000

Data Cleaning Procedure

- [REFORMAT COLUMN NAME](#)
- [FURTHER EXAMINE SOME COLUMNS: age, handicap, sms_received](#)
- [RENAME NO-SHOW TO SHOWED_UP AND INVERT VALUES](#)
- [CHANGE TO APPROPRIATE DATA TYPES](#)
- [REFORMAT SCHEDULED DAY and APPOINTMENT DAY](#)

REFORMAT COLUMN NAME

- replace '-' with '_'

- insert an underscore in between PatientID, AppointmentID, ScheduledDay and AppointmentDay
- change all letter to lower case
- rename ("Handcap" to handicap, "Hipertension" to "hypertension")

the code below will print a list of the old column names

```
In [9]: (df.columns)
```

```
Out[9]: Index(['PatientID', 'AppointmentID', 'Gender', 'ScheduledDay',
              'AppointmentDay', 'Age', 'Neighbourhood', 'Scholarship', 'Hipertension',
              'Diabetes', 'Alcoholism', 'Handcap', 'SMS_received', 'No-show'],
             dtype='object')
```

the code below will replace '-' with '_' (observe 'No-show')

```
In [10]: df.rename(columns=lambda x: x.strip().replace("-", "_"), inplace=True)
df.columns
```

```
Out[10]: Index(['PatientID', 'AppointmentID', 'Gender', 'ScheduledDay',
               'AppointmentDay', 'Age', 'Neighbourhood', 'Scholarship', 'Hipertension',
               'Diabetes', 'Alcoholism', 'Handcap', 'SMS_received', 'No_show'],
              dtype='object')
```

the regular expression below will insert an underscore to PatientId, AppointmentID, ScheduledDay and AppointmentDay and take all letters to lower case

```
In [11]: df.rename(columns=lambda x: re.sub(r'(?<!^)(?=[A-Z])', '_', x).lower(), inplace=True)
df.columns
```

```
Out[11]: Index(['patient_id', 'appointment_i_d', 'gender', 'scheduled_day',
               'appointment_day', 'age', 'neighbourhood', 'scholarship',
               'hipertension', 'diabetes', 'alcoholism', 'handcap', 's_m_s_received',
               'no_show'],
              dtype='object')
```

the code below will correct the effect of the regular expression by replacing (s_m_s with sms) (and i_d with id)

```
In [12]: df.rename(columns=lambda x: x.replace('s_m_s', 'sms'), inplace=True)
df.rename(columns=lambda x: x.replace('i_d', 'id'), inplace=True)
df.columns
```

```
Out[12]: Index(['patient_id', 'appointment_id', 'gender', 'scheduled_day',
               'appointment_day', 'age', 'neighbourhood', 'scholarship',
               'hipertension', 'diabetes', 'alcoholism', 'handcap', 'sms_received',
               'no_show'],
              dtype='object')
```

Finally, let's rename ("handcap" to handicap, "hipertension" to "hypertension")

```
In [13]: df.rename(columns={'handcap': 'handicap', 'hipertension': 'hypertension'}, inplace=True)
df.columns
```

```
Out[13]: Index(['patient_id', 'appointment_id', 'gender', 'scheduled_day',
               'appointment_day', 'age', 'neighbourhood', 'scholarship',
               'hypertension', 'diabetes', 'alcoholism', 'handicap', 'sms_received',
               'no_show'],
              dtype='object')
```

let's view the formatted column names below

```
In [14]: df.head()
```

```
Out[14]:
```

	patient_id	appointment_id	gender	scheduled_day	appointment_day	age	neighbourhood	s
0	2.987250e+13	5642903	F	2016-04-29T18:38:08Z	2016-04-29T00:00:00Z	62	JARDIM DA PENHA	
1	5.589978e+14	5642503	M	2016-04-29T16:08:27Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	
2	4.262962e+12	5642549	F	2016-04-29T16:19:04Z	2016-04-29T00:00:00Z	62	MATA DA PRAIA	
3	8.679512e+11	5642828	F	2016-04-29T17:29:31Z	2016-04-29T00:00:00Z	8	PONTAL DE CAMBURI	
4	8.841186e+12	5642494	F	2016-04-29T16:07:23Z	2016-04-29T00:00:00Z	56	JARDIM DA PENHA	

FURTHER EXAMINE SOME COLUMNS

from the summary statistics, we can see that

- age column have min value -1 and max value of 115years
- handicap column have max value 4 which is not a characteristics of a boolean
- sms_received column have min value of 0 and max value of 1 which is a characteristic of a boolean

we can conclude that:

- the age -1 would have been an error entry so we will we drop the rows having negative values of age
- the handicap and sms_received column would have been mistakenly swapped

let's check how many rows have a negative value in the age column

```
In [15]: print(sum(df.age < 0))
```

```
1
```

drop the single row having the negative value

```
In [16]: df.drop(df[df.age < 0].index, inplace=True)
```

HANDICAP Vs SMS_RECEIVED: first, compare the unique values in handicap and sms_received below

```
In [17]: print(f'"sms_received": {df.sms_received.unique()}')
print(f'"handicap": {df.handicap.unique()}')
```

```
"sms_received": [0 1]
"handicap":      [0 1 2 3 4]
```

let's swap the columns by renaming them below

```
In [18]: df.rename(columns={'sms_received': 'handicap', 'handicap': 'sms_received'}, inplace=True)
```

```
In [19]: print(f'"sms_received": {df.sms_received.unique()}')
print(f'"handicap": {df.handicap.unique()}')
```

```
"sms_received": [0 1 2 3 4]
"handicap":      [0 1]
```

RENAME NO_SHOW TO SHOWED_UP AND INVERT VALUES

in the original dataset, it says 'No' if the patient showed up to their appointment, and 'Yes' if they did not show up. this is because the column name is "no-show". We will change the column name to showed_up and invert the values in that column. i.e change 'No' to 'Yes' and vice versa

we must first convert no_show column to a boolean type so that we can perform the inversion operation

```
In [20]: df.no_show = df.no_show.apply(lambda x: x == 'Yes')
```

```
In [21]: df.dtypes
```

```
Out[21]: patient_id      float64
appointment_id    int64
gender            object
scheduled_day     object
appointment_day   object
age              int64
neighbourhood     object
scholarship       int64
hypertension      int64
diabetes          int64
alcoholism        int64
sms_received      int64
handicap          int64
no_show          bool
dtype: object
```

Now that 'no_show' is boolean, let's rename 'no_show' column and invert the values

```
In [22]: df['showed_up'] = ~df.no_show
```

let's confirm the inversion below

```
In [23]: df[['no_show', 'showed_up']].head()
```

```
Out[23]:
```

	no_show	showed_up
0	False	True
1	False	True
2	False	True
3	False	True
4	False	True

drop the no_show column

```
In [24]: df.drop('no_show', axis=1, inplace=True)
```

CHANGE TO APPROPRIATE DATA TYPES

from the data description, (**scholarship, hypertension, diabetes, alcoholism and handicap**) are supposed to be boolean, so we will convert them to boolean


```
In [25]: df['scholarship'] = df['scholarship'].astype(bool)
df['hypertension'] = df['hypertension'].astype(bool)
df['diabetes'] = df['diabetes'].astype(bool)
df['alcoholism'] = df['alcoholism'].astype(bool)
df['handicap'] = df['handicap'].astype(bool)
```

patient_id and **appointment_id** should be a string and not integer or floating number because they are used for identification. first, we will convert **patient_id** to integer so as to take care of decimal place and then convert both of them to strings

```
In [26]: df.patient_id = df.patient_id.astype('int64').astype(str)
df.appointment_id = df.appointment_id.astype(str)
```

REFORMAT SCHEDULED DAY and APPOINTMENT DAY

here, we will change the format of the date in the scheduled day and appointment day column to datetime

```
In [27]: df.scheduled_day = pd.to_datetime(df.scheduled_day, infer_datetime_format=True)
df.appointment_day = pd.to_datetime(df.appointment_day, infer_datetime_format=True)

df[['scheduled_day', 'appointment_day']].head()
```

Out[27]:

	scheduled_day	appointment_day
0	2016-04-29 18:38:08	2016-04-29
1	2016-04-29 16:08:27	2016-04-29
2	2016-04-29 16:19:04	2016-04-29
3	2016-04-29 17:29:31	2016-04-29
4	2016-04-29 16:07:23	2016-04-29

All datatype has been appropriately fixed.

In [28]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 110526 entries, 0 to 110526
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  -
0   patient_id            110526 non-null  object
1   appointment_id        110526 non-null  object
2   gender                 110526 non-null  object
3   scheduled_day          110526 non-null  datetime64[ns]
4   appointment_day        110526 non-null  datetime64[ns]
5   age                   110526 non-null  int64
6   neighbourhood          110526 non-null  object
7   scholarship            110526 non-null  bool
8   hypertension           110526 non-null  bool
9   diabetes               110526 non-null  bool
10  alcoholism             110526 non-null  bool
11  sms_received           110526 non-null  int64
12  handicap               110526 non-null  bool
13  showed_up              110526 non-null  bool
dtypes: bool(6), datetime64[ns](2), int64(2), object(4)
memory usage: 8.2+ MB
```

Cleaned Data

let us save the cleaned data below so that we can use it for the Exploratory Data Analysis

In [29]: `df.to_csv('C:/Users/user/Desktop/1.coding_datascience/UDACITY NANO DEGREE data an`

Exploratory Data Analysis

with our cleaned data, we are ready to explore. we will create visualizations with the goal of addressing the research questions that we posed in the Introduction section.

Load the cleaned data and check first few rows and columns.

```
In [30]: df = pd.read_csv('C:/Users/user/Desktop/1.coding_datascience/UDACITY NANO DEGREE')
df.head()
```

Out[30]:

	patient_id	appointment_id	gender	scheduled_day	appointment_day	age	neighbourhood
0	29872499824296	5642903	F	2016-04-29 18:38:08	2016-04-29 00:00:00	62	JARDIM DA PENHA
1	558997776694438	5642503	M	2016-04-29 16:08:27	2016-04-29 00:00:00	56	JARDIM DA PENHA
2	4262962299951	5642549	F	2016-04-29 16:19:04	2016-04-29 00:00:00	62	MATA DA PRAIA
3	867951213174	5642828	F	2016-04-29 17:29:31	2016-04-29 00:00:00	8	PONTAL DE CAMBUR
4	8841186448183	5642494	F	2016-04-29 16:07:23	2016-04-29 00:00:00	56	JARDIM DA PENHA


Research Question 1: Does Gender affect the rate of show up?

To know if gender affects the rate of show up, we have to know the total no of Male and Female patients and know the percentage of male and female who showed up for their appointment

check total number of Male and Female record.

```
In [31]: print(df.gender.value_counts())  
df.gender.hist();
```

```
F    71839  
M    38687  
Name: gender, dtype: int64
```



let's check the total no of Male and Female who showed up for their appointment

```
In [32]: gender_showed_up = df.groupby('gender').showed_up.value_counts()  
gender_showed_up
```

```
Out[32]: gender  showed_up  
F          True      57245  
          False     14594  
M          True     30962  
          False      7725  
Name: showed_up, dtype: int64
```

From the above counts:

57245 out of **71839** female patients showed_up for their appointment

30962 out of **38687** male patients showed_up for their appointment

let us compare their percentage in the chart below

```
In [33]: plt.figure(0)

# Data to plot
label1 = 'Females who showed up', 'Females who did not show up'
data1 = [57245, 14594]
color1 = ['gold', 'lightskyblue']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data1, explode=explode, labels=label1, colors=color1, autopct='%1.1f%%',

plt.axis('equal')


plt.figure(1)


# Data to plot
label2 = 'Males who showed up', 'Males who did not show up'
data2 = [30962, 7725]
color2 = ['yellowgreen', 'lightcoral']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data2, explode=explode, labels=label2, colors=color2, autopct='%1.1f%%',

plt.axis('equal')

plt.show()
```






from the above chart, we can see that Gender did not affect the rate of show up. roughly 80% of both gender showed up while roughly 20% of both gender did not show up

Research Question 2: Does the Bosnia Familia Scholarship affect the rate of show up?

let's check how many patients have scholarship

```
In [34]: print(df.scholarship.value_counts())
df.scholarship.astype(str).hist();
```

```
False    99665
True     10861
Name: scholarship, dtype: int64
```



10861 patients have the scholarship.

let us now check for the percentage of patients with scholarship who showed up for their appointment

```
In [35]: Bosnia = df.groupby('scholarship').showed_up.value_counts()
Bosnia
```

```
Out[35]: scholarship  showed_up
False              True      79924
              False      19741
True              True       8283
              False       2578
Name: showed_up, dtype: int64
```

From the above counts:

79924 out of 99665 regular patients showed_up for their appointment

8283 out of 10861 scholarship patients showed_up for their appointment

note that regular patients are those that don't have the bosnia scholarship

let us compare their percentage in the chart below

```
In [36]: plt.figure(0)

# Data to plot
label1 = 'Regular who showed up', 'Regular who did not show up'
data1 = [79924, 19741]
color1 = ['gold', 'lightskyblue']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data1, explode=explode, labels=label1, colors=color1, autopct='%1.1f%%',

plt.axis('equal')


plt.figure(1)

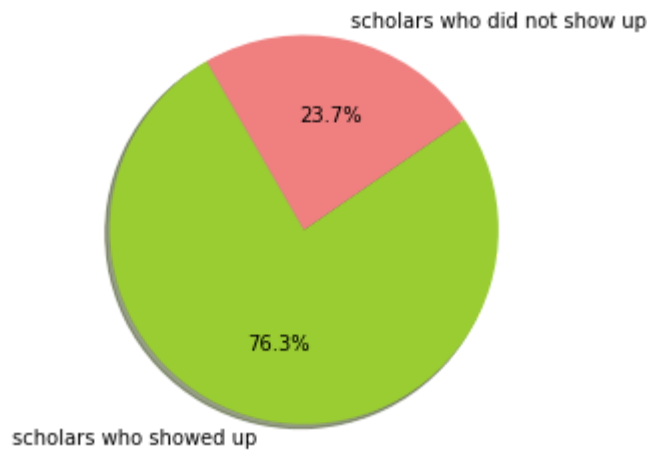
# Data to plot
label2 = 'scholars who showed up', 'scholars who did not show up'
data2 = [8283, 2578]
color2 = ['yellowgreen', 'lightcoral']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data2, explode=explode, labels=label2, colors=color2, autopct='%1.1f%%',

plt.axis('equal')

plt.show()
```






the Bosnia Familia Scholarship has an effect on rate of show up. Patients who do not have scholarship tend to show up more for their appointment.

Research Question 3: What is the effect of disability on show up?

```
In [37]: print(df.handicap.value_counts())
df.handicap.astype(str).hist();
```

```
False    75044
True     35482
Name: handicap, dtype: int64
```



```
In [38]: disable = df.groupby('handicap').showed_up.value_counts()
disable
```

```
Out[38]: handicap  showed_up
False      True         62509
           False        12535
True       True         25698
           False         9784
Name: showed_up, dtype: int64
```

From the above counts:

62509 out of 75044 patients without disability showed_up for their appointment

25698 out of 35482 patients with disability showed_up for their appointment

let us compare their percentage in the chart below

```
In [39]: plt.figure(0)

# Data to plot
label1 = 'able show', 'able no show'
data1 = [62509, 12535]
color1 = ['gold', 'lightskyblue']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data1, explode=explode, labels=label1, colors=color1, autopct='%1.1f%%',

plt.axis('equal')

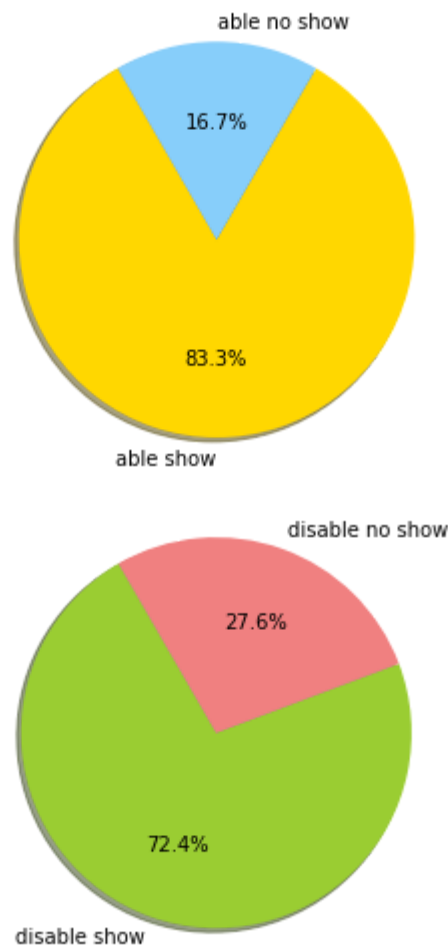
plt.figure(1)

# Data to plot
label2 = 'disable show', 'disable no show'
data2 = [25698, 9784]
color2 = ['yellowgreen', 'lightcoral']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data2, explode=explode, labels=label2, colors=color2, autopct='%1.1f%%',

plt.axis('equal')

plt.show()
```



disability affects the rate of show up, patients without disabilities shows up more for their appointment

Research Question 4: how was the bosnia scholarship distributed among patients?

```
In [40]: handischolar = df.groupby('handicap').scholarship.value_counts()  
handischolar
```

```
Out[40]: handicap  scholarship  
False      False      67688  
           True       7356  
True       False      31977  
           True       3505  
Name: scholarship, dtype: int64
```

```
In [41]: plt.figure(0)

# Data to plot
label1 = 'able sholar', 'able regular'
data1 = [7356, 67688]
color1 = ['gold', 'lightskyblue']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data1, explode=explode, labels=label1, colors=color1, autopct='%1.1f%%',

plt.axis('equal')

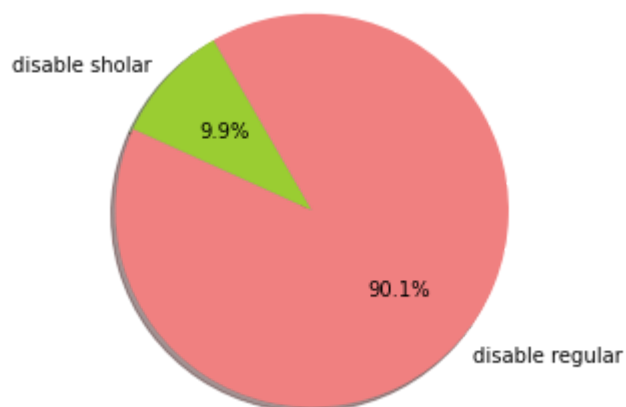
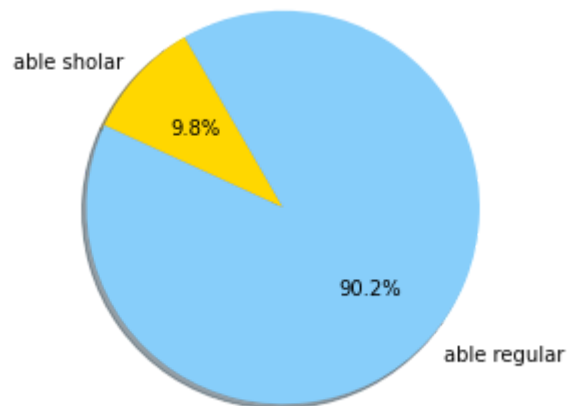
plt.figure(1)

# Data to plot
label2 = 'disable sholar', 'disable regular'
data2 = [3505, 31977]
color2 = ['yellowgreen', 'lightcoral']
explode = (0.0, 0.0) # explode 1st slice

# Plot
plt.pie(data2, explode=explode, labels=label2, colors=color2, autopct='%1.1f%%',

plt.axis('equal')

plt.show()
```



the bosnia scholarship was proportionately distributed between patients without disability(9.8%) and patients with disability(9.9%)

Conclusions

- females schedule more medical appointment than males. it suggest women takes way more care of their health in comparison to man
- the bosnia scholarship was proportionately distributed among patients with and without disability
- it seems the bosnia scholarship contributes to no show of patients
- comparing both genders who scheduled appointments, both genders have roughly the same show up rate

further Research Question: which of the features is the most important factor that determines if a patient showed-up.

In []: